

10_Amazon_Food_Reviews_Clustering

April 14, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 K-Means, Agglomerative & DBSCAN Clustering

Apply K-means Clustering on these feature sets:

-

SET 1: Review text, preprocessed one converted into vectors using (B

SET 2: Review text, preprocessed one converted into vectors

SET 3: Review text, preprocessed one converted into vectors

- SET 4:Review text, preprocessed one converted into vectorsFind the best k using the elbow-knee method (plot k vs inertia_)Once after you find the k clusters, plot the word cloud per each cluster so that at a single

go we can analyze the words in a cluster.

-

 Apply Agglomerative Clustering on these feature sets:

 SET 3:Review text, preprocessed one converted into vectors
 SET 4:Review text, preprocessed one converted into vectors
 Apply agglomerative algorithm and try a different number of clusters like 2,5 etc.
 Same as that of K-means, plot word clouds for each cluster and summarize in your own words w
 You can take around 5000 reviews or so(as this is very computationally expensive one

 Apply DBSCAN Clustering on these feature sets:

 SET 3:Review text, preprocessed one converted into vectors
 SET 4:Review text, preprocessed one converted into vectors
 Find the best Eps using the <a href='https://stackoverflow.com/questions/12893492/choosing-e
 Same as before, plot word clouds for each cluster and summarize in your own words what that
 You can take around 5000 reviews for this as well.

2.1 [A] K-Means Clustering

3 Import Required Packages

```
In [1]: import os
        from datetime import datetime
        import pandas as pd
        import numpy as np

        # package for reading from database
        import sqlite3

        # visualization related packages
        import matplotlib.pyplot as plt
        import seaborn as sns
        sns.set()

        # data prerocessing related
```

```

from sklearn.preprocessing import StandardScaler

# for computing distances between data points
from sklearn.metrics import pairwise_distances

# import model related packages
from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN
from sklearn.metrics import silhouette_score

# visualization related packages
from wordcloud import WordCloud
from prettytable import PrettyTable

```

4 UTIL Functions

4.1 Data Preprocessing related Functions

```

In [2]: def preprocess_data(config_dict, scaling=True, dim_reduction=False):
        """
        This function does preprocessing of data such as column standardization and
        dimensionality reduction using Truncated SVD
        """

        # Read train, test data frames & truncate it as needed
        train_df = pd.read_csv(config_dict['train_csv_path'], index_col=False)
        train_df = train_df.iloc[0:config_dict['train_size']]

        # print the statistics of train, test df
        print('Train df shape', train_df.shape)

        # separate features and labels
        id_values = train_df['Id']
        train_features = train_df.drop(['Label', 'Id'], axis=1)

        # get feature names as list
        feature_name_list = train_features.columns.values.tolist()

        # If Scaling is opted scale the train, test data
        if scaling:
            standard_scaler = StandardScaler()
            standard_scaler.fit(train_features)
            # scale the features
            train_features = pd.DataFrame(standard_scaler.transform(train_features),
                                          columns=feature_name_list)

        train_features['Id'] = id_values
        print('Shape of -> train features :%d,%d '%train_features.shape)

```

```

# if dim reduction is opted, reduce the dimension
if dim_reduction:
    # create an SVD object
    trunc_svd = TruncatedSVD(n_components=train_features.shape[1]-1, n_iter=8, algorithm='randomized')

    # fit to data
    trunc_svd.fit(train_features)
    # get explained variance ratio of each component
    explained_var_ratios = trunc_svd.explained_variance_ratio_
    # get cumulative ratio list for selecting the number of components
    cumulative_ratios = np.cumsum(explained_var_ratios)

    # plot the #components vs captured variance in the data
    plt.title('SVD Decomposition')
    plt.xlabel('Number of components')
    plt.ylabel('Cumulative Percentage Ratio')
    plt.plot(range(1, train_features.shape[1]), cumulative_ratios)
    plt.show()

    # set a threshold for stopping selection of components.
    svd_thesh = 0.001
    # select the number of components as the first component for which the difference
    # very less (less than svd thresh) compared with the very next component
    selected_dim = list(filter(lambda x : x[1] < svd_thesh, enumerate(np.diff(cumulative_ratios))))
    print('Num dimensions selected by SVD', selected_dim)
    print('Total variance captured:%f'%(cumulative_ratios[selected_dim]))

    # create an object for selecting the components
    trunc_svd = TruncatedSVD(n_components=selected_dim, n_iter=8, algorithm='randomized')
    # refit with the desired number of components
    trunc_svd.fit(train_features)

    # reduce the number of dimensions to selected number of components
    train_features = pd.DataFrame(trunc_svd.transform(train_features))

    # get the shape of final data frame and print it
    print('Shape of train df: (%d,%d)'%train_features.shape)

return train_features

```

4.2 Model training and evaluation related functions

In [3]: `def find_best_hyperparameter(config_dict, train_features):`

`"""`

This function helps to find the best hyper parameter for the clustering algorithm.

All set of hyper param values using which the model to be evaluated can be passed to list hyperparam_list. It also takes algo_type which can be one among k-means, agglom or DBSCAN.

"""

```
print('='*100)
```

```
hyperparam_list = config_dict['hyperparam_list']  
algo_type = config_dict['algo_type']
```

```
# set the input for training
```

```
X_train = train_features.drop(['Id'], axis=1)
```

```
hyper_param_score_list = list()
```

```
inertia_score_list = list() # for k-means algorithm
```

```
for hyp_vals in hyperparam_list:
```

```
# Model defined here
```

```
if algo_type == 'kmeans':
```

```
# create an object of clustering algorithm
```

```
clustering_model = KMeans(n_clusters=hyp_vals, init='k-means++', n_init=10)
```

```
# fit on data
```

```
clustering_model.fit(X_train)
```

```
# get the inertia value as score
```

```
inertia_val = clustering_model.inertia_
```

```
# get the silhouytte score for this clustering
```

```
if len(set(clustering_model.labels_)) == 1:
```

```
    sil_score = -1
```

```
else:
```

```
    sil_score = silhouette_score(X_train, clustering_model.labels_)
```

```
# append hyper param scores
```

```
inertia_score_list.append((hyp_vals, inertia_val,))
```

```
hyper_param_score_list.append((hyp_vals, sil_score,))
```

```
elif algo_type == 'agglomerative':
```

```
# create clustering algorithm object
```

```
clustering_model = AgglomerativeClustering(n_clusters=hyp_vals)
```

```
# fit on data
```

```
clustering_model.fit(X_train)
```

```
# get the silhouytte score for this clustering
```

```
if len(set(clustering_model.labels_)) == 1:
```

```

        sil_score = -1
    else:
        sil_score = silhouette_score(X_train, clustering_model.labels_)

    sil_score = silhouette_score(X_train, clustering_model.labels_)

    # append hyper param scores
    hyper_param_score_list.append((hyp_vals, sil_score,))

elif algo_type == 'dbscan':

    # hyper params in the order (eps, min_samples,)
    clustering_model = DBSCAN(eps=hyp_vals[0], min_samples=hyp_vals[0], metric='
    # fit on data
    clustering_model.fit(X_train)
    # get the inertia value as score

    if len(set(clustering_model.labels_)) == 1:
        sil_score = -1 # minimum possible silhoutte score

    else:
        sil_score = silhouette_score(X_train, clustering_model.labels_)

    # append hyper param scores
    hyper_param_score_list.append((hyp_vals, sil_score,))

else:
    print('Invalid choice')

# plot inertia vs k for k-means clustering
if algo_type == 'kmeans':
    # plot k versus inertia
    inertia_values_list = [item[1] for item in inertia_score_list]
    plt.plot(hyperparam_list, inertia_values_list)
    plt.xlabel('Number of Clusters')
    plt.ylabel('Inertia Value')
    plt.title('Elbow Method')
    plt.show()

return hyper_param_score_list

```

```
In [4]: def get_cluster_id(config_dict, score_list, train_features):
```

```

    """
    This function fit a model based on best hyper parameter values got and
    assign cluster id to each data point.
    """

```

```

# get the best hyperparams from the list
hyp_vals, hyp_score = max(score_list, key=lambda x: x[1])

# get configuration values
hyperparam_list = config_dict['hyperparam_list']
algo_type = config_dict['algo_type']

# Model defined here
if algo_type == 'kmeans':
    print('Best hyper param selected n_clusters : %d '%hyp_vals)
    print('Best silhoutte score for this hyper parameter : %f'%hyp_score)
    clustering_model = KMeans(n_clusters=hyp_vals, init='k-means++', n_init=10)
    # set entry for putting in table
    hyp_str = 'k= ' + str(hyp_vals)

elif algo_type == 'agglomerative':
    print('Best hyper param selected n_clusters : %d'%hyp_vals)
    print('Best silhoutte score score for this hyper parameter : %f'%hyp_score)
    clustering_model = AgglomerativeClustering(n_clusters=hyp_vals)
    # set entry for putting in table
    hyp_str = 'k= ' + str(hyp_vals)

elif algo_type == 'dbscan':
    # hyper params in the order (eps, min_samples,)
    print('Best hyper param selected eps:%f,min_samples :%d'%hyp_vals)
    print('Best silhoutte score score for this hyper parameter : %f'%hyp_score)
    clustering_model = DBSCAN(eps=hyp_vals[0], min_samples=hyp_vals[1], metric='eucl

    # set entry for putting in table
    hyp_str = 'eps= %.4f, min_pts= %d '%hyp_vals

else:
    print('Invalid choice')

X_train = train_features.drop(['Id'], axis=1)

clustering_model.fit(X_train)

# get the labels for each data point (i.e cluster id)
assigned_cluseter_ids = clustering_model.labels_

# create a data frame with review ID and its predicted cluster
pred_df = pd.DataFrame({'Id': train_features['Id'],
                        'Cluster' : assigned_cluseter_ids},
                        index=range(train_features.shape[0]))

# set the clustering size info as a string
cluster_info = str(dict(pred_df['Cluster'].value_counts()))

```

```

# form a table entry to insert into pretty table
table_entry = [cluster_info, hyp_str, '{0:.4f}'.format(hyp_score)]

return (table_entry, pred_df,)

```

```

In [5]: def get_kdistance_plot(df, k_val):
        """
        This function returns a k-distance plot for a given data frame. It compute all pairwise
        and then compute average k-distances. It also finds the eps_val for DBSCAN algorithm
        """

        # compute pairwise distances and create a distance data frame
        pair_dist_df = pd.DataFrame(pairwise_distances(df))
        pair_dist_df.shape[0]

        # declare a list for holding mean k-distances
        mean_k_distance_list = list()

        # for each data point compute the k-mean distance
        for index, row in pair_dist_df.iterrows():

            # drop the distance to the point itself
            temp_entry = row.drop([index])

            # sort the values in ascending order of distance
            temp_entry = temp_entry.sort_values(ascending=True)
            # pick the first k_val points as the k distance
            temp_entry = temp_entry.iloc[0:k_val]

            # compute the mean k distance
            mean_k_distance = temp_entry.mean()

            # save the k mean distance for this point into a list
            mean_k_distance_list.append(mean_k_distance)

        # sort the k-mean distance in descending order
        mean_k_distance_list = sorted(mean_k_distance_list, reverse=True)

        # find eps value by looking into the sudden change in distance
        maxima_diff_list = list(filter(lambda x : x > 0.005, abs(np.diff(mean_k_distance_list))))
        eps_val = mean_k_distance_list[len(maxima_diff_list)]

        # plot the k-mean distances in a figure
        plt.plot(range(1, pair_dist_df.shape[0] + 1), mean_k_distance_list)
        plt.xlabel('Data Points')
        plt.ylabel('K distance')

```



```
plt.title('K distance plot -> k:%d, eps:%f'%(k_val, eps_val,))
plt.show()

return (mean_k_distance_list, eps_val,)
```

4.3 Functions for preparing word cloud of each cluster

```
In [6]: def read_from_DB(id_list):
        """
        This function reads the Cleaned Review Text and its id into a data frame and returns
        """

        # create a connection object for connecting with DB
        con = sqlite3.connect('/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/fina

        #
        df = pd.read_sql_query('SELECT Id, CleanedText, Summary, Label from Reviews', con)

        # get only the selected id and the correspondig CleanedText
        df = df[df['Id'].isin(id_list)]
        df = df[['Id', 'CleanedText', 'Label']]

        # close the connection
        con.close()

        return df

In [7]: def get_cluster_wordcloud(pred_cluster_df, review_id_df):
        """
        This function shows the word cloud for each cluster given the predicted cluster info
        """

        # create a word cloud template for each cluster word cloud
        wc = WordCloud(background_color='white', width=800, height=800)

        # do word cloud for each cluster
        for gid, gdf in pred_cluster_df.groupby(['Cluster']):

            print('='*100)

            # get id of each data point in this cluster
            id_list = gdf['Id'].tolist()

            # filter only the selected id reviews
            selected_rev_df = review_id_df[review_id_df['Id'].isin(id_list)]
            # get the class (+ve, -ve) of each ID
            val_counts_dict = dict(selected_rev_df['Label'].value_counts())
```

```

# get the majority class for this cluster
num_pos_rev = val_counts_dict[1]
num_neg_rev = val_counts_dict[0]
majority_class = 'Positive' if num_pos_rev > num_neg_rev else 'Negative'

# print a table containing information about this cluster
Pret_table_c = PrettyTable()
Pret_table_c.field_names = ['Cluster_ID', 'Num Data Points', '# +ve Reviews',
                             '# -ve Reviews', 'Majority Class']
Pret_table_c.title = 'Cluster %d Info'%(gid,)

Pret_table_c.add_row([gid, gdf.shape[0], val_counts_dict[1], val_counts_dict[0],
                      majority_class])
print(Pret_table_c)

# initialize a text for this cluster, this will hold all the words that belong to
cluster_text = str()

for rev_txt in selected_rev_df['CleanedText']:
    cluster_text += rev_txt

# generate word cloud for this cluster
wc_output = wc.generate(cluster_text)

plt.figure(figsize=(8,8))
plt.imshow(wc_output)
plt.axis('off')
plt.title('Word Cloud for Cluster ID: %d'%(gid,))
plt.tight_layout(pad=0.0)
plt.show()

print('='*100)

```

4.3.1 [A.1] Applying K-Means Clustering on BOW, SET 1

```

In [8]: config_dict = {
        'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIEW/BOW/train_data.csv',
        'train_size' : 15000,
        'hyperparam_list' : [2, 3, 4, 5, 6],
        'algo_type' : 'kmeans'
    }

```


```

In [9]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)
        score_list = find_best_hyperparameter(config_dict, train_features)

```

Train df shape (15000, 503)

Shape of -> train features :15000,502



```
In [10]: print('Score list for this clustering :\n', score_list)
```

Score list for this clustering :

```
[(2, 0.35220317632255554), (3, 0.20534017832980486), (4, 0.1498655129908616), (5, 0.14238491192)
```

```
In [11]: ptabe_entry_a1, pred_df = get_cluster_id(config_dict, score_list, train_features)
         pred_df.head()
```

Best hyper param selected n_clusters : 2

Best silhoutte score for this hyper parameter : 0.352203

```
Out[11]:
```

	Id	Cluster
0	456873	1
1	81416	0
2	519340	0
3	340949	0
4	453782	0

4.3.2 [A.2] Wordclouds of clusters obtained after applying k-means on BOW SET 1

```
In [12]: # get id of each review in the training data points
         id_list = train_features['Id'].tolist()
         # get only the reviews which are used for training step
         rev_df = read_from_DB(id_list)
         # plot the word cloud for every clusters predicted
         get_cluster_wordcloud(pred_df, rev_df)
```

```
=====
+-----+-----+-----+-----+-----+
|                                     Cluster 0 Info                                     |
+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+
|      0      |      13051      |      6551      |      6500      |      Positive   |
+-----+-----+-----+-----+-----+
```


[illegible]

Inertia value decreases sharply without any elbow region

Both clusters have many words in common. This means its not the word but a pattern of word that distinguishes the cluster. We cannot simply relay on just one feature to do clustering

4.4.1 [A.3] Applying K-Means Clustering on TFIDE, SET 2


```
In [13]: config_dict = {  
        'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/TFIDF/t  
        'train_size' : 15000,  
        'hyperparam_list' : [2, 3, 4, 5, 6],  
        'algo_type' : 'kmeans'  
    }
```

```
In [14]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)  
        score_list = find_best_hyperparameter(config_dict, train_features)
```

Train df shape (15000, 503)

Shape of -> train features :15000,502

=====



```
In [15]: print('Score list for this clustering :\n', score_list)
```

Score list for this clustering :

[(2, 0.005332590945988098), (3, 0.0035891003118323508), (4, 0.0048517343200489595), (5, -0.010000000000000001), (6, -0.010000000000000001)]

```
In [16]: ptabe_entry_a2, pred_df = get_cluster_id(config_dict, score_list, train_features)  
        pred_df.head()
```

Best hyper param selected n_clusters : 2
Best silhoutte score for this hyper parameter : 0.005333

```
Out[16]:
```

	Id	Cluster
0	456873	0
1	81416	0
2	519340	1
3	340949	1
4	453782	0

4.4.2 [A.4] Wordclouds of clusters obtained after applying k-means on TFIDF SET 2

```
In [17]: # get id of each review in the training data points
id_list = train_features['Id'].tolist()
# get only the reviews which are used for training step
rev_df = read_from_DB(id_list)
# plot the word cloud for every clusters predicted
get_cluster_wordcloud(pred_df, rev_df)
```

```
=====
+-----+-----+-----+-----+-----+
|                                     Cluster 0 Info                                     |
+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+
|      0      |      6780      |      2923      |      3857      |      Negative      |
+-----+-----+-----+-----+-----+
```


Word Cloud for Cluster ID: 0



Cluster 1 Info				
Cluster_ID	Num Data Points	# +ve Reviews	# -ve Reviews	Majority Class
1	8220	4501	3719	Positive

[illegible]

Inertia value decreases sharply without any elbow region.

Both clusters have many words in common. This means it is not the word but a pattern of word that distinguishes the cluster. We cannot simply rely on just one feature to do clustering

18


4.5.1 [A.5] Applying K-Means Clustering on AVG W2V, SET 3

```
In [18]: config_dict = {  
        'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/AVG_W2V',  
        'train_size' : 15000,  
        'hyperparam_list' : [2, 3, 4, 5, 6],  
        'algo_type' : 'kmeans'  
    }
```

```
In [19]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)  
        score_list = find_best_hyperparameter(config_dict, train_features)
```

Train df shape (15000, 52)

Shape of -> train features :15000,51



```
In [20]: print('Score list for this clustering :\n', score_list)
```

Score list for this clustering :

[(2, 0.09443577923088471), (3, 0.09690401266701287), (4, 0.09556962073504416), (5, 0.0850925880...)]

```
In [21]: ptabe_entry_a3, pred_df = get_cluster_id(config_dict, score_list, train_features)  
        pred_df.head()
```

Best hyper param selected n_clusters : 3
Best silhoutte score for this hyper parameter : 0.096904

```
Out[21]:
```

	Id	Cluster
0	456873	1
1	81416	2
2	519340	0
3	340949	1
4	453782	1

4.5.2 [A.6] Wordclouds of clusters obtained after applying k-means on AVG W2V SET 3

```
In [22]: # get id of each review in the training data points
id_list = train_features['Id'].tolist()
# get only the reviews which are used for training step
rev_df = read_from_DB(id_list)
# plot the word cloud for every clusters predicted
get_cluster_wordcloud(pred_df, rev_df)
```

```
=====
+-----+-----+-----+-----+-----+
|                                     Cluster 0 Info                                     |
+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+
|      0      |      2393      |      1309      |      1084      |      Positive      |
+-----+-----+-----+-----+-----+
```


[illegible]

Inertia value decreases sharply without any elbow region

Both clusters have many words in common. This means it's not the word but a pattern of words that distinguishes the cluster. We cannot simply rely on just one feature to do clustering


4.6.1 [A.7] Applying K-Means Clustering on TFIDF W2V, SET 4

```
In [23]: config_dict = {  
        'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/TFIDF_W  
        'train_size' : 15000,  
        'hyperparam_list' : [2, 3, 4, 5, 6],  
        'algo_type' : 'kmeans'  
    }
```

```
In [24]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)  
        score_list = find_best_hyperparameter(config_dict, train_features)
```

Train df shape (15000, 52)

Shape of -> train features :15000,51



```
In [25]: print('Score list for this clustering :\n', score_list)
```

Score list for this clustering :

[(2, 0.15845545839009997), (3, 0.09276123777834097), (4, 0.09837859926179662), (5, 0.0996582982

```
In [26]: ptabe_entry_a4, pred_df = get_cluster_id(config_dict, score_list, train_features)  
        pred_df.head()
```


Best hyper param selected n_clusters : 2
Best silhoutte score for this hyper parameter : 0.158455

```
Out[26]:
```

	Id	Cluster
0	456873	0
1	81416	0
2	519340	1
3	340949	0
4	453782	0

4.6.2 [A.8] Wordclouds of clusters obtained after applying k-means on TFIDF W2V SET 4

```
In [27]: # get id of each review in the training data points
id_list = train_features['Id'].tolist()
# get only the reviews which are used for training step
rev_df = read_from_DB(id_list)
# plot the word cloud for every clusters predicted
get_cluster_wordcloud(pred_df, rev_df)
```

```
=====
+-----+-----+-----+-----+-----+
|                                     Cluster 0 Info                                     |
+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+
|      0      |      12946      |      6322      |      6624      |      Negative   |
+-----+-----+-----+-----+-----+
```


[illegible]

Inertia value decreases sharply without any elbow region

Both clusters have many words in common. This means it is not the word but a pattern of word that distinguishes the cluster. We cannot simply rely on just one feature to do clustering

4.8 [B] Agglomerative Clustering

4.8.1 [B.1] Applying Agglomerative Clustering on AVG W2V, SET 3

```
In [28]: config_dict = {
          'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/AVG_W2V',
          'train_size' : 5000,
          'hyperparam_list' : [2, 3, 4, 5, 6],
          'algo_type' : 'agglomerative'
        }
```

```
In [29]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)
          score_list = find_best_hyperparameter(config_dict, train_features)
```

Train df shape (5000, 52)

Shape of -> train features :5000,51

```
In [30]: print('Score list for this clustering :\n', score_list)
```

Score list for this clustering :

[(2, 0.06843433509933095), (3, 0.06813809061058529), (4, 0.06849660399005532), (5, 0.0483884027

```
In [31]: ptabe_entry_b1, pred_df = get_cluster_id(config_dict, score_list, train_features)
          pred_df.head()
```

Best hyper param selected n_clusters : 4

Best silhoutte score score for this hyper parameter : 0.068497

```
Out[31]:
```

	Id	Cluster
0	456873	1
1	81416	1
2	519340	2
3	340949	3
4	453782	0


4.8.2 [B.2] Wordclouds of clusters obtained after applying Agglomerative Clustering on AVG W2V SET 3

```
In [32]: # get id of each review in the training data points
          id_list = train_features['Id'].tolist()
          # get only the reviews which are used for training step
          rev_df = read_from_DB(id_list)
          # plot the word cloud for every clusters predicted
          get_cluster_wordcloud(pred_df, rev_df)
```

Cluster 0 Info				
Cluster_ID	Num Data Points	# +ve Reviews	# -ve Reviews	Majority Class
0	1573	722	851	Negative



Cluster 1 Info				
Cluster_ID	Num Data Points	# +ve Reviews	# -ve Reviews	Majority Class
1	2177	1150	1027	Positive



```

+-----+
|               Cluster 2 Info               |
+-----+

```



```

        'algo_type' : 'agglomerative'
    }

In [34]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)
        score_list = find_best_hyperparameter(config_dict, train_features)

Train df shape (5000, 52)
Shape of -> train features :5000,51
=====

In [35]: print('Score list for this clustering :\n', score_list)

Score list for this clustering :
[(2, 0.13184396652056307), (3, 0.1176629658795971), (4, 0.09243798682890239), (5, 0.09142444357

In [36]: ptabe_entry_b2, pred_df = get_cluster_id(config_dict, score_list, train_features)
        pred_df.head()

Best hyper param selected n_clusters : 2
Best silhoutte score score for this hyper parameter : 0.131844

Out[36]:
      Id Cluster
0  456873      0
1   81416      0
2  519340      1
3  340949      1
4  453782      0

```

4.8.4 [B.4] Wordclouds of clusters obtained after applying Agglomerative Clustering on TFIDF W2V SET 4

```

In [37]: # get id of each review in the training data points
        id_list = train_features['Id'].tolist()
        # get only the reviews which are used for training step
        rev_df = read_from_DB(id_list)
        # plot the word cloud for every clusters predicted
        get_cluster_wordcloud(pred_df, rev_df)

```

```

=====
+-----+-----+-----+-----+-----+-----+
|                                     Cluster 0 Info                                     |
+-----+-----+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+-----+-----+
|      0      |      4383      |      2174      |      2209      |      Negative  |
+-----+-----+-----+-----+-----+-----+

```


[illegible]

Silhouette coefficient can be used to select the number of clusters

35

4.10 [C] DBSCAN Clustering

4.10.1 [C.1] Applying DBSCAN on AVG W2V, SET 3

```
In [38]: config_dict = {  
    'train_csv_path' : '/home/amd_3/AAIC/ASM_REPO/Processed_data/AMZN_FOOD_REVIW/AVG_W2V',  
    'train_size' : 5000,  
    'min_pts_list' : [40, 60, 80, 100, 120, 150],  
    'hyperparam_list' : list(), # initialize this list as empty  
    'algo_type' : 'dbscan'  
}
```

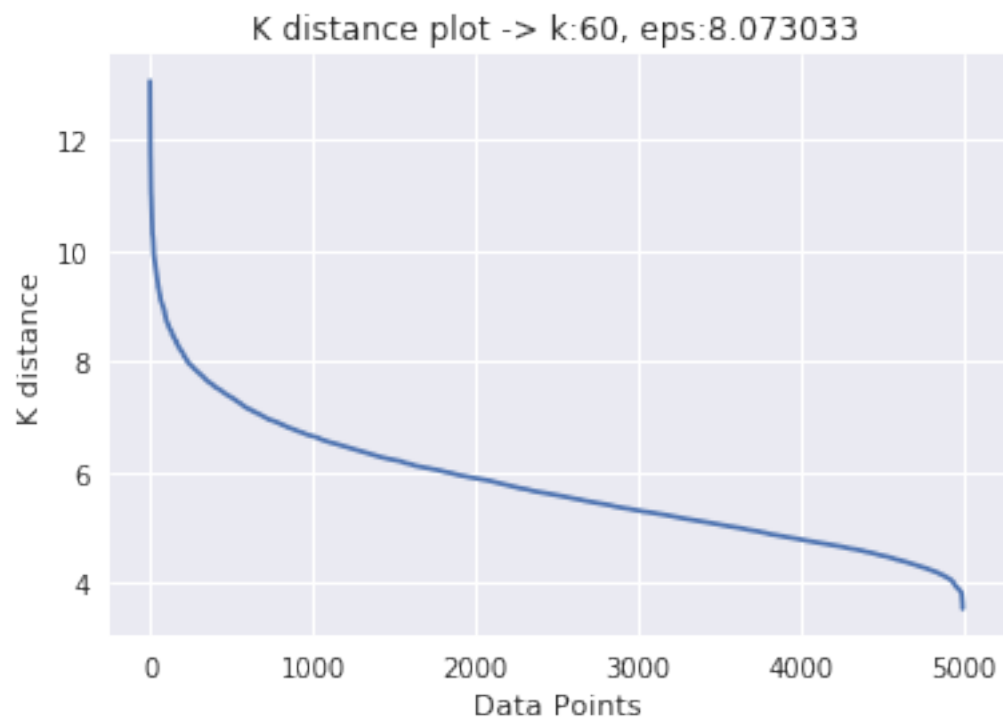

```
In [39]: train_features = preprocess_data(config_dict, scaling=True, dim_reduction=False)
```



Train df shape (5000, 52)



Shape of -> train features :5000,51

4.10.2 Identify the eps_val using k-distance method

```
In [40]: dist_input_df = train_features.drop(['Id'], axis=1)  
  
    # declare a list to hold all the hyperparams  
    hyp_param_list = list()  
  
    for min_pts in config_dict['min_pts_list']:  
        mean_k_distance_list, eps_val, = get_kdistance_plot(dist_input_df, min_pts)  
        hyp_param_list.append((eps_val, min_pts,))  
  
    # set hyper param in  
    config_dict['hyperparam_list'] = hyp_param_list
```







```
In [41]: score_list = find_best_hyperparameter(config_dict, train_features)
        print('Score list for this clustering :\n', score_list)
```

```
=====
Score list for this clustering :
```

```
[(7.879124555167806, 40), 0.31914734436834824), ((8.073033236962353, 60), 0.3326226160676315),
```

```
In [42]: ptabe_entry_c1, pred_df = get_cluster_id(config_dict, score_list, train_features)
        pred_df.head()
```

```
Best hyper param selected eps:8.597458,min_samples :150
```

```
Best silhoutte score score for this hyper parameter : 0.355103
```

```
Out[42]:
```

	Id	Cluster
0	456873	0
1	81416	0
2	519340	0
3	340949	0
4	453782	0

4.10.3 [C.2] Wordclouds of clusters obtained after applying DBSCAN on AVG W2V SET 3

```
In [43]: # get id of each review in the training data points
        id_list = train_features['Id'].tolist()
        # get only the reviews which are used for training step
        rev_df = read_from_DB(id_list)
        # plot the word cloud for every clusters predicted
        get_cluster_wordcloud(pred_df, rev_df)
```


```
=====
+-----+-----+-----+-----+-----+-----+
|                                     Cluster -1 Info                                     |
+-----+-----+-----+-----+-----+-----+-----+
| Cluster_ID | Num Data Points | # +ve Reviews | # -ve Reviews | Majority Class |
+-----+-----+-----+-----+-----+-----+-----+
|      -1    |         21      |         10     |         11     |      Negative  |
+-----+-----+-----+-----+-----+-----+-----+
```







```
Train df shape (5000, 52)  
Shape of -> train features :5000,51
```

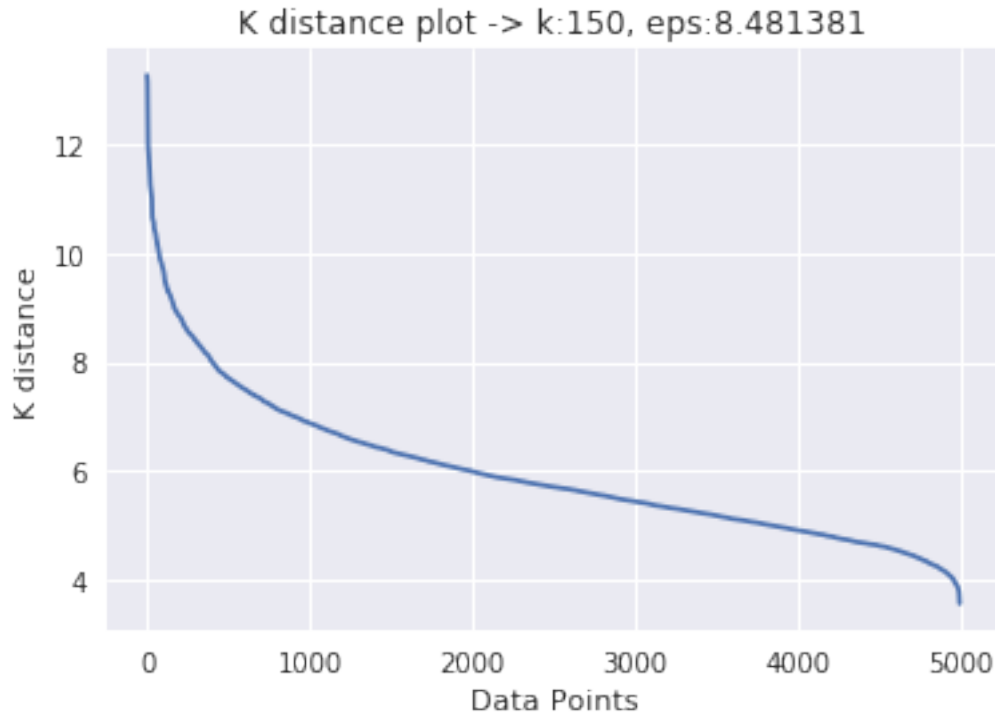
4.10.5 Identify the eps_val using k-distance method

```
In [46]: dist_input_df = train_features.drop(['Id'], axis=1)  
  
# declare a list to hold all the hyperparams  
hyp_param_list = list()  
  
for min_pts in config_dict['min_pts_list']:  
    mean_k_distance_list, eps_val, = get_kdistance_plot(dist_input_df, min_pts)  
    hyp_param_list.append((eps_val, min_pts,))  
  
# set hyper param in  
config_dict['hyperparam_list'] = hyp_param_list
```









```
In [47]: score_list = find_best_hyperparameter(config_dict, train_features)
         print('Score list for this clustering :\n', score_list)
```

=====

Score list for this clustering :

[((7.437505298428046, 40), 0.3000192553249008), ((7.82774585189994, 60), 0.3425005647524928), (

```
In [48]: ptabe_entry_c2, pred_df = get_cluster_id(config_dict, score_list, train_features)
         pred_df.head()
```

Best hyper param selected eps:8.287429,min_samples :120

Best silhoutte score score for this hyper parameter : 0.386138


```
Out[48]:
```

	Id	Cluster
0	456873	0
1	81416	0
2	519340	0
3	340949	0
4	453782	0

4.10.6 [C.4] Wordclouds of clusters obtained after applying DBSCAN on TFIDF W2V SET 4

```
In [49]: # get id of each review in the training data points
         id_list = train_features['Id'].tolist()
```


Cluster 0 Info				
Cluster_ID	Num Data Points	# +ve Reviews	# -ve Reviews	Majority Class
0	4988	2502	2486	Positive



4.11 Observation

Silhouette coefficient can be used to select the number of clusters

Cluster ID -1 (Noise Points) has many words which are related to groceries such as rice, coconut oil, tomato etc.

DBSCAN method identified one noisy cluster (-1) and one valid cluster (0)

5 Procedure Summary

All the four datasets are column standardized before feeding to clustering algorithm

Best hyper parameter is selected using inertia/silhouette score. In this assignment for all clustering methods the clusters are selected based on the silhouette score

All words in a cluster are displayed using word cloud representation

The min_pts points are set by trial and error method

For DBSCAN the right number of eps value is determined using k-distance plot. From k-distance plot we need to identify the knee point in order to set the eps value.

6 Results Summary

```
In [50]: Pret_table = PrettyTable()
Pret_table.field_names = ['Vectorizer', 'Algorithm', 'Cluster Info {id:size}', 'Hyper-P
Pret_table.title = 'Clustering Results Summary'
```

```
In [51]: # K-means
Pret_table.add_row(['BoW', 'K-means'] + ptabe_entry_a1)
Pret_table.add_row(['TF-IDF', 'K-means'] + ptabe_entry_a2)
Pret_table.add_row(['Avg W2V', 'K-means'] + ptabe_entry_a3)
Pret_table.add_row(['TF-IDF W2V', 'K-means'] + ptabe_entry_a4)

# Agglomerative
Pret_table.add_row(['Avg W2V', 'Agglomerative'] + ptabe_entry_b1)
Pret_table.add_row(['TF-IDF W2V', 'Agglomerative'] + ptabe_entry_b2)

# DBSCAN
Pret_table.add_row(['Avg W2V', 'DBSCAN'] + ptabe_entry_c1)
Pret_table.add_row(['TF-IDF W2V', 'DBSCAN'] + ptabe_entry_c2)
```

```
In [52]: print(Pret_table)
```

```
+-----+
|                                     Clustering Results Summary                                     |
+-----+-----+-----+-----+-----+
| Vectorizer | Algorithm | Cluster Info {id:size} | Hyper-Param |
+-----+-----+-----+-----+-----+
| BoW        | K-means  | {0: 13051, 1: 1949}   | k= 2        |
| TF-IDF     | K-means  | {1: 8220, 0: 6780}    | k= 2        |
| Avg W2V    | K-means  | {2: 6776, 1: 5831, 0: 2393} | k= 3        |
| TF-IDF W2V | K-means  | {0: 12946, 1: 2054}   | k= 2        |
```

Avg W2V	Agglomerative	{1: 2177, 0: 1573, 3: 633, 2: 617}	k= 4	
TF-IDF W2V	Agglomerative	{0: 4383, 1: 617}	k= 2	
Avg W2V	DBSCAN	{0: 4979, -1: 21}	eps= 8.5975, min_pts= 150	
TF-IDF W2V	DBSCAN	{0: 4988, -1: 12}	eps= 8.2874, min_pts= 120	
+-----+-----+-----+-----+				

7 Conclusions

The best clustering obtained is from DBSCAN with silhouette value 0.3861

Many clusters formed have many words in common. This means that we cannot simply relay on words to get proper clustering, rather we need to try with many combination of words .