

Quora_Model

June 29, 2019

```
In [1]: # general purpose packages
import pandas as pd
import numpy as np

import os
from datetime import datetime
import pickle

# visualization related packages
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import train_test_split

# generate random distributions for hyperparams
from scipy.stats import randint
from scipy.stats import beta
from scipy.stats import gamma
from scipy.stats import expon
from scipy.stats import uniform
from scipy.stats import poisson

# Classifier Evaluation
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_recall_fscore_support
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import log_loss
from sklearn.metrics import make_scorer

# data preprocessing
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import normalize
```

```

# import model related packages
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
import xgboost
from xgboost import XGBClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.model_selection import StratifiedKFold

from wordcloud import WordCloud
from sklearn.manifold import TSNE

from prettytable import PrettyTable

```

1 Configs

```

In [2]: sample_size = -1 # set -1 if we want to use full size # set fractional value for sample

train_df_path = './data/Final_train_df.csv'
test_df_path = './data/Final_test_df.csv'

```

2 Util Functions

```

In [3]: def get_confusion_matrix(actual_list, predicted_list, title_suffix=str()):
        """
        This function plots the confusion matrix given ground truth and predicted
        """

        conf_matrix = confusion_matrix(actual_list, predicted_list, labels=[0, 1])

        col_names = ['Non Duplicate', 'Duplicate']
        conf_df = pd.DataFrame(conf_matrix, columns=col_names)
        conf_df.index = col_names

        # plot confusion matrix
        sns.heatmap(conf_df, annot=True, annot_kws={'size': 16}, fmt='g', cmap='YlGnBu',
                    cbar_kws={'label': 'prediction_count', 'format': '%d'})
        plt.xticks(rotation=0)
        plt.yticks(rotation=90)
        plt.xlabel('Predicted')
        plt.ylabel('Actual')
        plt.title('Confusion Matrix -' + title_suffix)
        plt.show()

```

```
return conf_df
```

```
In [4]: def get_precision_recall_matrix(conf_matrix, title_suffix=str()):
```

```
    # compute precision matrix
```

```
    precision_matrix = conf_matrix.div(conf_matrix.sum(axis=0), axis=1) * 100.0
```

```
    # compute recall matrix
```

```
    recall_matrix = conf_matrix.div(conf_matrix.sum(axis=1), axis=0) * 100.0
```

```
    # plot both the matrices
```

```
    # plot precision matrix
```

```
    sns.heatmap(precision_matrix, annot=True, annot_kws={'size': 16}, fmt='.4f', cmap='YlGn',  
                cbar_kws={'label': 'percentage', 'format': '%.2f'})
```

```
    plt.yticks(rotation=0)
```

```
    plt.xticks(rotation=90)
```

```
    plt.xlabel('Predicted')
```

```
    plt.ylabel('Actual')
```

```
    plt.title('Precision Matrix -' + title_suffix)
```

```
    plt.show()
```

```
    # plot recall matrix
```

```
    sns.heatmap(recall_matrix, annot=True, annot_kws={'size': 16}, fmt='.4f', cmap='YlGn',  
                cbar_kws={'label': 'percentage', 'format': '%.2f'})
```

```
    plt.yticks(rotation=0)
```

```
    plt.xticks(rotation=90)
```

```
    plt.xlabel('Predicted')
```

```
    plt.ylabel('Actual')
```

```
    plt.title('Recall Matrix -' + title_suffix)
```

```
    plt.show()
```

```
    # return as a tuple
```

```
    return (precision_matrix, recall_matrix,)
```

```
In [5]: def get_classification_report(actual, predicted, title_suffix=str()):
```

```
    # set class labels and its corresponding name
```

```
    class_labels_list = [0, 1]
```

```
    col_names = ['Non Duplicate', 'Duplicate']
```

```
    # compute performance df
```

```
    eval_matrix = precision_recall_fscore_support(actual, predicted,  
                                                  labels=class_labels_list)
```

```
    eval_df = pd.DataFrame(list(eval_matrix), columns=class_labels_list)
```

```
    eval_df.index = ['Precision', 'Recall', 'Fscore', 'Support']
```

```

eval_df.columns = col_names

# normalize the performace df
eval_df_normed = eval_df * 100.0
eval_df_normed.loc['Support', col_names] /= eval_df_normed.loc['Support', col_names]
eval_df_normed.iloc[3:4, :] *= 100.0

# plot the classification report
sns.heatmap(eval_df_normed, annot=True, annot_kws={'size': 16}, fmt='.4f', cmap='YlGn',
            cbar_kws={'label': 'Percentage', 'format': '%.2f'})
plt.yticks(rotation=0)
plt.xticks(rotation=90)

plt.xlabel('Classes')
plt.ylabel('Metrics')
plt.title('Classification Report -' + title_suffix)
plt.show()

return eval_df_normed

```

```

In [6]: def show_class_distribution(df, title):
        """
        This function displays the distribution of each class as a barchart
        """
        count_df = pd.DataFrame(list(dict(df['is_duplicate'].value_counts()).items()),
                                columns=['is_duplicate', 'Count'])

        plt.title('Class Distribution of ' + title + ' Data')
        plt.xlabel('is_duplicate')
        plt.ylabel('Count')
        sns.barplot(x='is_duplicate', y=count_df['Count'], data=count_df)

        plt.show()

```

```

In [7]: def evaluate_model(model, X, y):

        # get prediction and its probability
        predicted_labels = model.predict(X)
        predicted_probs = model.predict_proba(X)

        # get confsuion matrix
        conf_matrix = get_confusion_matrix(y, predicted_labels)

        # get precision & recall matrix
        pre_matrix, recall_matrix = get_precision_recall_matrix(conf_matrix)

        # get classification report

```

```

clf_report = get_classification_report(y, predicted_labels)

# get the logistic loss
log_loss_value = log_loss(y, predicted_probs, labels=[0, 1], eps=1e-15)
log_loss_value = round(log_loss_value, 4)

table_entry = (log_loss_value,)

return table_entry

```

3 Data

```

In [8]: df_train = pd.read_csv(train_df_path, index_col=False)
df_test = pd.read_csv(test_df_path, index_col=False)

```

```

print('Shape of train data : ', df_train.shape)
print('Shape of test data : ', df_test.shape)

if sample_size > 0:
    print('Sample is taken ', sample_size)
    # train sample
    df_train = df_train.sample(frac=sample_size)
    df_train = df_train.reset_index(drop=True)
    # test sample
    df_test = df_test.sample(frac=sample_size)
    df_test = df_test.reset_index(drop=True)

    print('Shape of train data : ', df_train.shape)
    print('Shape of test data : ', df_test.shape)

```

```

Shape of train data : (84000, 123)
Shape of test data : (36000, 123)

```

```

In [9]: df_train.head()

```

```

Out[9]:

```

	id	is_duplicate	q1len	q2len	q1_n_words	q2_n_words	word_Common	\
0	223477	0	73	61	12	13	2	
1	111372	0	32	40	6	6	3	
2	291123	0	67	46	13	10	2	
3	317134	1	33	37	7	6	4	
4	246990	0	94	58	16	10	1	

	word_Total	word_share	ctc_min	...	q2_start	q2_think	q2_time	q2_use	\
0	23	0.086957	0.999999	...	0.0	0.0	0.0	0.0	
1	12	0.250000	0.999998	...	0.0	0.0	0.0	0.0	
2	22	0.090909	0.923076	...	0.0	0.0	0.0	0.0	
3	13	0.307692	0.999999	...	0.0	0.0	0.0	0.0	

4	25	0.040000	0.937499	...	0.0	0.0	0.0	0.0
---	----	----------	----------	-----	-----	-----	-----	-----

	q2_want	q2_way	q2_without	q2_work	q2_would	q2_year
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 123 columns]

```
In [10]: X_train = df_train.drop(['id', 'is_duplicate'], axis=1)
y_train = df_train['is_duplicate']

X_test = df_test.drop(['id', 'is_duplicate'], axis=1)
y_test = df_test['is_duplicate']
```

3.1 Scale the data

```
In [11]: col_names = X_train.columns.values.tolist()
```

```
# fit to data
std_scaler = StandardScaler()
std_scaler.fit(X_train)

# scale the datasets
X_train = pd.DataFrame(std_scaler.transform(X_train), columns=col_names)
X_test = pd.DataFrame(std_scaler.transform(X_test), columns=col_names)

X_train.head()
```

```
/home/nisheel-s/anaconda3/lib/python3.6/site-packages/sklearn/preprocessing/data.py:625: DataConversionWarning:
  return self.partial_fit(X, y)
/home/nisheel-s/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: DataConversionWarning:
/home/nisheel-s/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: DataConversionWarning:
  if __name__ == '__main__':
```

```
Out[11]:
```

	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	\
0	0.449286	0.027934	0.195881	0.288942	-0.811405	0.222577	
1	-0.921050	-0.589063	-0.910307	-0.812989	-0.487533	-1.063037	
2	0.248749	-0.412778	0.380245	-0.183314	-0.811405	0.105703	
3	-0.887627	-0.677205	-0.725942	-0.812989	-0.163660	-0.946163	
4	1.151165	-0.060208	0.933339	-0.183314	-1.135278	0.456325	

	word_share	ctc_min	ctc_max	cwc_min	...	q2_start	q2_think	\
0	-1.061825	0.537775	0.537775	0.013822	...	-0.109526	-0.119748	
1	0.207553	0.537762	0.537762	0.013725	...	-0.109526	-0.119748	

```

2  -1.031052 -0.637190 -0.637190  0.013876 ... -0.109526 -0.119748
3   0.656718  0.537765  0.537765  0.013725 ... -0.109526 -0.119748
4  -1.427406 -0.416882 -0.416882  0.013846 ... -0.109526 -0.119748

```

```

      q2_time    q2_use    q2_want    q2_way  q2_without  q2_work  q2_would  \
0 -0.133215 -0.123439 -0.108397 -0.148691  -0.105842 -0.116595 -0.156094
1 -0.133215 -0.123439 -0.108397 -0.148691  -0.105842 -0.116595 -0.156094
2 -0.133215 -0.123439 -0.108397 -0.148691  -0.105842 -0.116595 -0.156094
3 -0.133215 -0.123439 -0.108397 -0.148691  -0.105842 -0.116595 -0.156094
4 -0.133215 -0.123439 -0.108397 -0.148691  -0.105842 -0.116595 -0.156094

```

```

      q2_year
0 -0.102217
1 -0.102217
2 -0.102217
3 -0.102217
4 -0.102217

```

```
[5 rows x 121 columns]
```

4 Model

```
In [12]: def predict_using_random_model(df, df_type):
```

```
    # create a random prediction dataframe
```

```
    rand_df = pd.DataFrame(np.random.rand(df.shape[0], 2))
```

```
    # convert each row into a probability distribution
```

```
    rand_df = rand_df.div(rand_df.sum(axis=1), axis=0)
```

```
    # compute log loss
```

```
    log_loss_val = log_loss(df['is_duplicate'], rand_df.values, labels=[0,1], eps=1e-15)
```

```
    # get log loss on validation from the random model prediction
```

```
    print('Log loss on %s data using Random Model : %f'%(df_type, log_loss_val,))
```

```
    #return rand_df
```

```
In [13]: predict_using_random_model(df_train, 'Train')
```

```
         predict_using_random_model(df_test, 'Train')
```

```
Log loss on Train data using Random Model : 0.883145
```

```
Log loss on Train data using Random Model : 0.884132
```

4.1 1. Logistic Regression

```
In [14]: def get_best_hyperparam_LogisticRegression(param_dict, X, y, random_search=False):

    print(datetime.now() , ' Hyper param tuning of logistic regression started')

    # set the scoring function
    final_scorer = 'log_loss'

    # declare a scoring dictionary
    score_dict = {
        'log_loss': make_scorer(score_func=log_loss, greater_is_better=False,
                                needs_proba=True, needs_threshold=False,
                                eps=1e-15, labels=[0,1])
    }

    #Declare the metric as 'minimization' or 'maximization'
    optimization_dict = {
        'log_loss' : 'minimization'
    }

    # set data partitioning method
    cv_data = 3

    # declare model
    model = LogisticRegression()

    if random_search:
        search_cv = RandomizedSearchCV(estimator=model, param_distributions=param_dict,
                                       cv=cv_data, scoring=score_dict, refit=False,
                                       return_train_score=True, n_iter=6, n_jobs=-1)
    else:
        # declare grid search CV object
        search_cv = GridSearchCV(estimator=model, param_grid=param_dict, cv=cv_data,
                                scoring=score_dict, refit=False,
                                return_train_score=True, n_jobs=-1)

    # fit to the data
    search_cv.fit(X, y)

    # get total number of param settings
    param_list = list(param_dict.keys())
    param_field_list = ['param_' + item for item in param_list]
```



```

# get list of train metric list
train_metric_list = ['mean_train_' + item for item in score_dict.keys()]
# get list of test metric list
test_metric_list = ['mean_test_' + item for item in score_dict.keys()]

# get number of rows in the search cv data frame
num_rows = len(search_cv.cv_results_['params'])

# create the grid search info df
grid_info_df = pd.DataFrame(search_cv.cv_results_, index=range(num_rows))

# prepare a list of required columns
required_columns = ['params'] + param_field_list + train_metric_list + \
    test_metric_list

# slice the data frame to only required columns
grid_info_df = grid_info_df[required_columns]

# process individual metrics
for metric, optimization in optimization_dict.items():
    if optimization == 'minimization':
        grid_info_df['mean_train_' + metric] *= -1
        grid_info_df['mean_test_' + metric] *= -1

# Find the best hyperparam & its corresponding scores
if optimization_dict[final_scorer] == 'minimization':
    best_hyperparam_record = grid_info_df.loc[grid_info_df[
        'mean_test_' + final_scorer].idxmin(),:]
else:
    best_hyperparam_record = grid_info_df.loc[grid_info_df[
        'mean_test_' + final_scorer].idxmax(),:]

# best hyperparam & corresponding scores
best_hyperparam = best_hyperparam_record['params']
best_train_score = best_hyperparam_record['mean_train_' + final_scorer]
best_validation_score = best_hyperparam_record['mean_test_' + final_scorer]

# plot the hyper params
if len(param_list) == 1:

    # extract individual fields
    x_vals = np.log10(grid_info_df[param_field_list[0]].tolist())
    y_vals_tr = grid_info_df['mean_train_' + final_scorer].tolist()
    y_vals_val = grid_info_df['mean_test_' + final_scorer].tolist()

    # plot train, validation performances
    plt.plot(x_vals, y_vals_tr, label='Train', color='r', marker='o', linestyle='-')

```

```

plt.plot(x_vals, y_vals_val, label='Validation', color='b', marker='o', linestyle='--')
plt.xlabel('Log Hyper Params -' + param_list[0])
plt.ylabel(final_scorer)
plt.legend()
plt.title('LogisticRegression - Hyperparam Train v/s Validation Scores')
plt.show()

# Heatmap plot for pair of hyperparam values
elif len(param_list) == 2:
    # get pivoted table
    train_hyp_df = pd.pivot_table(data=grid_info_df, index=param_field_list[0],
                                   columns=param_field_list[1],
                                   values='mean_train_' + final_scorer,
                                   fill_value=np.inf)

    sns.heatmap(train_hyp_df, annot=True, cmap='YlGnBu', fmt='.4f',
                 cbar_kws={'label': final_scorer, 'format': '%.2f'})
    plt.xlabel(param_list[1])
    plt.ylabel(param_list[0])
    plt.title('LogisticRegression - Hyperparams Scores - Train')
    plt.show()

    print('\n'*3)

    # Test hyperparam

    # get pivoted table
    train_hyp_df = pd.pivot_table(data=grid_info_df, index=param_field_list[0],
                                   columns=param_field_list[1],
                                   values='mean_test_' + final_scorer, fill_value=np.inf)

    sns.heatmap(train_hyp_df, annot=True, cmap='YlGnBu', fmt='.4f',
                 cbar_kws={'label': final_scorer, 'format': '%.2f'})
    plt.xlabel(param_list[1])
    plt.ylabel(param_list[0])
    plt.title('LogisticRegression - Hyperparams Scores - Validation')
    plt.show()

else:
    print(grid_info_df)

print('Best hyperparam value: ', best_hyperparam, 'Best Train Score: ', best_train_score,
      'Best Validation Score: ', best_validation_score)

#best_mse_train = best_hyperparam_record['mean_train_MSE']
#best_mse_validation = best_hyperparam_record['mean_test_MSE']

# return tuple

```

```

ret_tuple = (best_hyperparam, best_train_score, best_validation_score,)

print(datetime.now() , ' Hyper param tuning of logistic regression Completed')

return ret_tuple

```

4.1.1 1. Find best hyperparameter


```

In [15]: # declare a set of params to search for
param_dict_lr = {'C' : [1e-04, 1e-03, 1e-02, 1e-01, 1e+00, 1e+01, 1e+02]}

hyp_tuned_info = get_best_hyperparam_LogisticRegression(param_dict_lr, X_train,
                                                         y_train, False)
best_hyp_lr, best_tr_score_lr, best_val_score_lr = hyp_tuned_info
lr_clf = LogisticRegression(C=best_hyp_lr['C'])

```

2019-06-26 14:36:52.765010 Hyper param tuning of logistic regression started



Best hyperparam value: {'C': 1.0} Best Train Score: 0.523874261103707 Best Validation Score: 0.526125738896293
 2019-06-26 14:37:51.933136 Hyper param tuning of logistic regression Completed

4.1.2 2. Train the model with best hyperparameter

```
In [16]: # train the model using the best hyperparam obtained
print(datetime.now() , ' Training of logistic regression started')
sig_lr_clf = CalibratedClassifierCV(base_estimator=lr_clf, method='sigmoid', cv=3)
sig_lr_clf.fit(X_train, y_train)
print(datetime.now() , ' Training of logistic regression Completed')

# save model to disk
pickle_out = open("./model/logistic_regression.pkl", "wb")
pickle.dump(sig_lr_clf, pickle_out)
pickle_out.close()
```

2019-06-26 14:37:51.941957 Training of logistic regression started


/home/nisheel-s/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning


/home/nisheel-s/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning


/home/nisheel-s/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning


2019-06-26 14:38:16.352678 Training of logistic regression Completed

```
In [17]: # get table entry for training data
#load model from disk
pickle_in = open("./model/logistic_regression.pkl", "rb")
sig_lr_clf = pickle.load(pickle_in)
pickle_in.close()
table_entry_lr_train = evaluate_model(sig_lr_clf, X_train, y_train)
```













4.1.3 3. Test the model

```
In [18]: table_entry_lr_test = evaluate_model(sig_lr_clf, X_test, y_test)
```





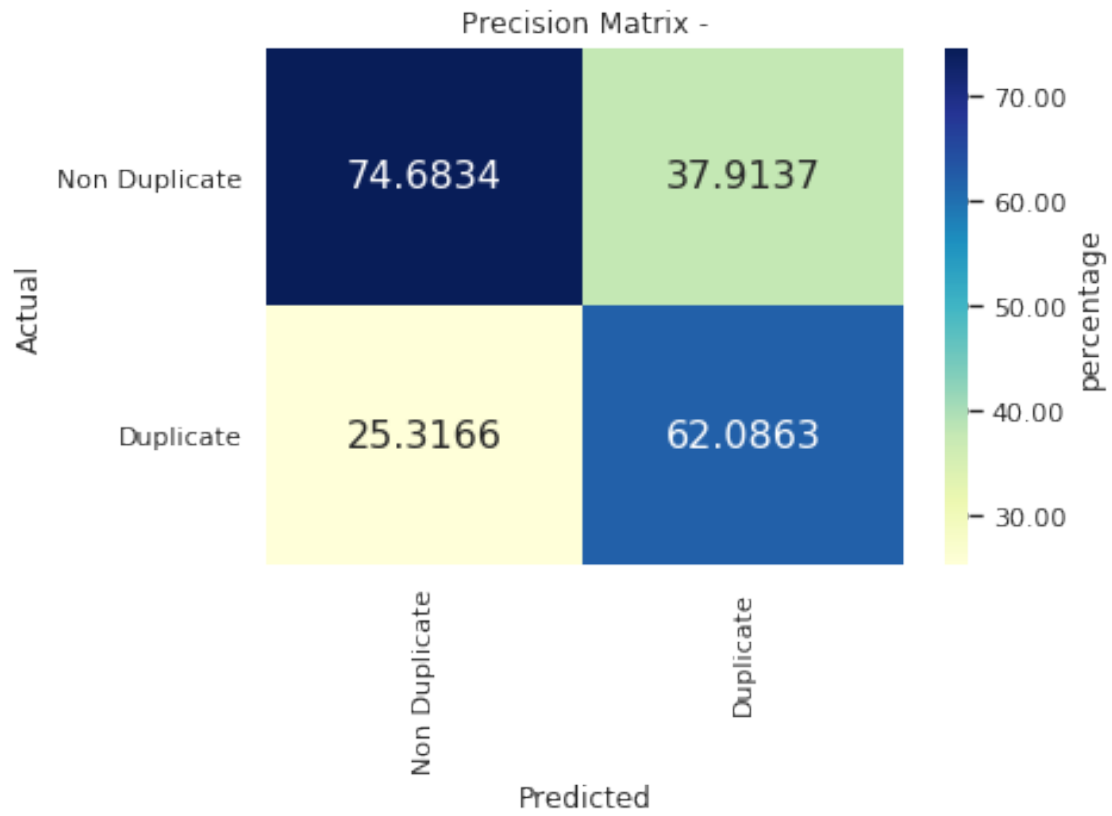








```
In [19]: table_entry_lr_test = evaluate_model(sig_lr_clf, X_test, y_test)
         table_entry_lr = ('Logistic Regression', best_hyp_lr,) + table_entry_lr_train + table_e
```









4.2 2. Linear SVM Classifier

```
In [20]: def get_best_hyperparam_SVM_Classifier(param_list, X, y):

    print(datetime.now() , ' Hyperparam tuning of SVM started')

    # set the scoring function
    final_scorer = 'log_loss'

    # set data partitioning method
    cv_data = 3

    # set stratified K fold validation
    skf = StratifiedKFold(n_splits=cv_data)

    #declare two list for holding the loss for train, validation
    tr_loss_all = list()
    val_loss_all = list()
```



```

for alpha_val in param_list:

    # declare model
    model = SGDClassifier(loss='hinge', alpha=alpha_val, tol=1e-03,
                           max_iter=1e+03, penalty='l2', n_jobs=-1)

    # declare calibrated classifier
    calib_model = CalibratedClassifierCV(base_estimator=model,
                                         method='sigmoid',
                                         cv='prefit')

    # declare two lists for train, validation loss for each fold
    train_loss_list = list()
    val_loss_list = list()

    # evaluate for each fold
    for train_index, val_index in skf.split(X, y):

        # get train test split for this fold
        X_train, X_val = X.loc[train_index,:], X.loc[val_index,:]
        y_train, y_val = y[train_index], y[val_index]

        # fit to data
        model.fit(X_train, y_train)
        calib_model.fit(X_val, y_val)

        # get train, test log loss
        tr_probs = calib_model.predict_proba(X_train)
        val_probs = calib_model.predict_proba(X_val)

        # insert into log loss list
        train_loss_list.append(log_loss(y_train, tr_probs,
                                         labels=[0,1], eps=1e-15))
        val_loss_list.append(log_loss(y_val, val_probs,
                                      labels=[0,1], eps=1e-15))

    # compute mean loss values
    mean_tr_loss = np.mean(train_loss_list)
    mean_val_loss = np.mean(val_loss_list)

    # update the lists
    tr_loss_all.append(mean_tr_loss)
    val_loss_all.append(mean_val_loss)

# get total number of param settings
# get number of rows in the search cv data frame
num_rows = len(param_list)

```

```

# create the grid search info df
cv_results_dict = {'alpha': param_list,
                   'mean_train_log_loss': tr_loss_all,
                   'mean_test_log_loss': val_loss_all}

grid_info_df = pd.DataFrame(cv_results_dict, index=range(num_rows))

# prepare a list of required columns
required_columns = ['alpha', 'mean_train_log_loss',
                   'mean_test_log_loss']

# slice the data frame to only required columns
grid_info_df = grid_info_df[required_columns]

# Find the best hyperparam & its corresponding scores
best_hyperparam_record = grid_info_df.loc[grid_info_df[
    'mean_test_' + final_scorer].idxmin(),:]

# best hyperparam & corresponding scores
best_hyperparam = {'alpha': best_hyperparam_record['alpha']}
best_train_score = best_hyperparam_record['mean_train_' + final_scorer]
best_validation_score = best_hyperparam_record['mean_test_' + final_scorer]

# plot the hyper params
# extract individual fields
x_vals = np.log10(param_list)
y_vals_tr = grid_info_df['mean_train_' + final_scorer].tolist()
y_vals_val = grid_info_df['mean_test_' + final_scorer].tolist()

# plot train, validation performances
plt.plot(x_vals, y_vals_tr, label='Train', color='r', marker='o', linestyle='-')
plt.plot(x_vals, y_vals_val, label='Validation', color='b', marker='o', linestyle='-')
plt.xlabel('Log Hyper Params - Alpha')
plt.ylabel(final_scorer)
plt.legend()
plt.title('SVM Classifier - Hyperparam Train v/s Validation Scores')
plt.show()

print('Best hyperparam value: ', best_hyperparam, 'Best Train Score: ', best_train_score,
      'Best Validation Score: ', best_validation_score)

# return tuple
ret_tuple = (best_hyperparam, best_train_score, best_validation_score,)

print(datetime.now(), ' Hyperparam tuning of SVM completed')

```

```
return ret_tuple
```

4.2.1 1. Find best hyperparameter


```
In [21]: # declare a set of params to search for
param_list_svm = [1e-05, 1e-04, 1e-03, 1e-02, 1e-01, 1e+00, 1e+01]

hyp_tuned_info = get_best_hyperparam_SVM_Classifier(param_list_svm, X_train,
                                                    y_train)
best_hyp_svm, best_tr_score_svm, best_val_score_svm = hyp_tuned_info

svm_clf = SGDClassifier(loss='hinge', alpha=best_hyp_svm['alpha'], tol=1e-03,
                        max_iter=1e+03, penalty='l2', n_jobs=-1)

svm_sig_clf = CalibratedClassifierCV(base_estimator=svm_clf,
                                     method='sigmoid', cv=3)
```

2019-06-26 14:38:18.633282 Hyperparam tuning of SVM started



Best hyperparam value: {'alpha': 0.01} Best Train Score: 0.5305024792875814 Best Validation Score: 0.533282
2019-06-26 14:38:52.919826 Hyperparam tuning of SVM completed

4.2.2 2. Train the model with best hyperparameter


```
In [22]: # train the model using the best hyperparam obtained
print(datetime.now() , ' Training of SVM started')
svm_sig_clf.fit(X_train, y_train)
print(datetime.now() , ' Training of SVM completed')


# save model to disk
pickle_out = open("./model/svm.pkl", "wb")
pickle.dump(svm_sig_clf, pickle_out)
pickle_out.close()
```


2019-06-26 14:38:52.930837 Training of SVM started


2019-06-26 14:38:54.522214 Training of SVM completed

```
In [23]: pickle_in = open("./model/svm.pkl", "rb")
svm_sig_clf = pickle.load(pickle_in)
pickle_in.close()
table_entry_svm_train = evaluate_model(svm_sig_clf, X_train, y_train)
```













4.2.3 3. Test the model

```
In [24]: table_entry_svm_test = evaluate_model(svm_sig_clf, X_test, y_test)
         table_entry_svm = ('Linear SVM', best_hyp_svm,) + table_entry_svm_train + table_entry_s
```









4.3 3. XG Boost Classifier

```
In [25]: def get_best_hyperparam_XGBClassifier(param_dict, X, y, random_search=False):

    print(datetime.now(), ' Hyperparam Tuning of XGB started')

    # set the scoring function
    final_scorer = 'log_loss'

    # declare a scoring dictionary
    score_dict = {
        'log_loss': make_scorer(score_func=log_loss, greater_is_better=False,
                                needs_proba=True, needs_threshold=False,
                                eps=1e-15, labels=[0,1])
    }

    #Declare the metric as 'minimization' or 'maximization'
    optimization_dict = {
```

```

        'log_loss' : 'minimization'
    }

    # set data partitioning method
    cv_data = 3

    # declare model
    model = XGBClassifier()

    if random_search:
        search_cv = RandomizedSearchCV(estimator=model, param_distributions=param_dict,
                                       cv=cv_data, scoring=score_dict, refit=False,
                                       return_train_score=True, n_iter=6, n_jobs=-1)

    else:
        # declare grid search CV object
        search_cv = GridSearchCV(estimator=model, param_grid=param_dict, cv=cv_data,
                                scoring=score_dict, refit=False,
                                return_train_score=True, n_jobs=-1)

    # fit to the data
    search_cv.fit(X, y)

    # get total number of param settings
    param_list = list(param_dict.keys())
    param_field_list = ['param_' + item for item in param_list]

    # get list of train metric list
    train_metric_list = ['mean_train_' + item for item in score_dict.keys()]
    # get list of test metric list
    test_metric_list = ['mean_test_' + item for item in score_dict.keys()]

    # get number of rows in the search cv data frame
    num_rows = len(search_cv.cv_results_['params'])

    # create the grid search info df
    grid_info_df = pd.DataFrame(search_cv.cv_results_, index=range(num_rows))

    # prepare a list of required columns
    required_columns = ['params'] + param_field_list + train_metric_list + \
        test_metric_list

    # slice the data frame to only required columns
    grid_info_df = grid_info_df[required_columns]

    # process individual metrics
    for metric, optimization in optimization_dict.items():

```

```

        if optimization == 'minimization':
            grid_info_df['mean_train_' + metric] *= -1
            grid_info_df['mean_test_' + metric] *= -1

# Find the best hyperparam & its corresponding scores
    if optimization_dict[final_scorer] == 'minimization':
        best_hyperparam_record = grid_info_df.loc[grid_info_df[
            'mean_test_' + final_scorer].idxmin(),:]
    else:
        best_hyperparam_record = grid_info_df.loc[grid_info_df[
            'mean_test_' + final_scorer].idxmax(),:]

# best hyperparam & corresponding scores
    best_hyperparam = best_hyperparam_record['params']
    best_train_score = best_hyperparam_record['mean_train_' + final_scorer]
    best_validation_score = best_hyperparam_record['mean_test_' + final_scorer]

# plot the hyper params
    if len(param_list) == 1:

        # extract individual fields
        x_vals = grid_info_df[param_field_list[0]].tolist()
        y_vals_tr = grid_info_df['mean_train_' + final_scorer].tolist()
        y_vals_val = grid_info_df['mean_test_' + final_scorer].tolist()

        # plot train, validation performances
        plt.plot(x_vals, y_vals_tr, label='Train', color='r', marker='o', linestyle='-')
        plt.plot(x_vals, y_vals_val, label='Validation', color='b', marker='o', linestyle='-')
        plt.xlabel('Hyper Params -' + param_list[0])
        plt.ylabel(final_scorer)
        plt.legend()
        plt.title('XGB Classifier - Hyperparam Train v/s Validation Scores')
        plt.show()

# Heatmap plot for pair of hyperparam values
    elif len(param_list) == 2:
        # get pivoted table
        train_hyp_df = pd.pivot_table(data=grid_info_df, index=param_field_list[0],
                                       columns=param_field_list[1],
                                       values='mean_train_' + final_scorer,
                                       fill_value=np.inf)
        sns.heatmap(train_hyp_df, annot=True, cmap='YlGnBu', fmt='.4f',
                    cbar_kws={'label': final_scorer, 'format': '%.2f'})
        plt.xlabel(param_list[1])
        plt.ylabel(param_list[0])
        plt.title('XGB Classifier - Hyperparams Scores - Train')
        plt.show()

```

```

print('\n'*3)

# Test hyperparam

# get pivoted table
train_hyp_df = pd.pivot_table(data=grid_info_df, index=param_field_list[0],
                              columns=param_field_list[1],
                              values='mean_test_' + final_scorer, fill_value=np.nan)
sns.heatmap(train_hyp_df, annot=True, cmap='YlGnBu', fmt='.4f',
            cbar_kws={'label': final_scorer, 'format': '%.2f'})
plt.xlabel(param_list[1])
plt.ylabel(param_list[0])
plt.title('XGB Classifier - Hyperparams Scores - Validation')
plt.show()

else:
    print(grid_info_df)

print('Best hyperparam value: ', best_hyperparam, 'Best Train Score: ', best_train_score,
      'Best Validation Score: ', best_validation_score)

#best_mse_train = best_hyperparam_record['mean_train_MSE']
#best_mse_validation = best_hyperparam_record['mean_test_MSE']

# return tuple
ret_tuple = (best_hyperparam, best_train_score, best_validation_score,)

print(datetime.now(), ' Hyperparam Tuning of XGB completed')

return ret_tuple

```

4.3.1 1. Find best hyperparameter

```

In [26]: # declare a set of params to search for
param_dict_xgb = {'n_estimators' : [20, 60, 120, 200, 250],
                  'max_depth' : [2, 5, 20, 100, 250]
                  }

hyp_tuned_info = get_best_hyperparam_XGBClassifier(param_dict_xgb, X_train,
                                                  y_train, False)
best_hyp_xgb, best_tr_score_xgb, best_val_score_xgb = hyp_tuned_info


xgb_clf = XGBClassifier(n_estimators=best_hyp_xgb['n_estimators'],
                        max_depth=best_hyp_xgb['max_depth'])


```

```
xgb_sig_clf = CalibratedClassifierCV(base_estimator=xgb_clf, method='sigmoid', cv=3)
```

2019-06-26 14:38:55.850678 Hyperparam Tuning of XGB started

/home/nisheel-s/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/pr
"timeout or by a memory leak.", UserWarning





Best hyperparam value: {'max_depth': 20, 'n_estimators': 60} Best Train Score: 0.2204272962825
 2019-06-26 15:58:46.269411 Hyperparam Tuning of XGB completed


4.3.2 2. Train the model with best hyperparameter


```
In [27]: # train the model using the best hyperparam obtained
print(datetime.now() , ' Training of XGB started')
xgb_sig_clf.fit(X_train, y_train)
print(datetime.now() , ' Training of XGB completed')


# save model to disk
pickle_out = open("./model/xgb.pkl","wb")
pickle.dump(xgb_sig_clf, pickle_out)
pickle_out.close()
```


2019-06-26 15:58:46.280307 Training of XGB started
 2019-06-26 16:01:21.236148 Training of XGB completed

```
In [28]: pickle_in = open("./model/xgb.pkl","rb")
xgb_sig_clf = pickle.load(pickle_in)
pickle_in.close()
table_entry_xgb_train = evaluate_model(xgb_sig_clf, X_train, y_train)
```











4.3.3 3. Test the model

```
In [29]: table_entry_xgb_test = evaluate_model(xgb_sig_clf, X_test, y_test)
         best_hyp_xgb = '\n'.join(str(best_hyp_xgb).split(','))
         table_entry_xgb = ('XG Boost', best_hyp_xgb,) + table_entry_xgb_train + table_entry_xgb
```









5 Results

```
In [30]: Pret_table = PrettyTable()
        Pret_table.field_names = ['Model', 'Hyperparam', 'Train_Log_Loss', 'Test_Log_Loss']
        Pret_table.title = 'Classification Model Results Summary'

        # enter model performance metrics
        Pret_table.add_row(table_entry_lr)
        Pret_table.add_row(table_entry_svm)
        Pret_table.add_row(table_entry_xgb)
        print(Pret_table)
```

Model	Hyperparam	Train_Log_Loss	Test_Log_Loss
Logistic Regression	{'C': 1.0}	0.5242	0.5313
Linear SVM	{'alpha': 0.01}	0.5301	0.5372
XG Boost	{'max_depth': 20 'n_estimators': 60}	0.2753	0.4493

+-----+-----+-----+-----+

6 Procedure Summary

Three ML models (LogisticRegression, LinearSVM, XGB) are used to solve the problem

- Hyperparam tuning of all models done

- All models are trained with the best hyperparam values

- The performacne of each model is evaluated on a test set

7 Conclusion

The best log loss obtained is for XGB model (0.4493)

The XGB model showed a tendency to overfit due to quiet large deviation of loss between train test data

- More feature engineering methods can be tried to imporove the results further

- The f1-score for XGB test is 69.19 % , and SVM - 56.70%, Logistic regression- 57.22%

- XGB model outperformed other models interms of f1-score