

06_Implement_SGD

April 14, 2019

1 Import Required Packages

```
In [1]: import pandas as pd
import numpy as np

# preprocessing related packages
from sklearn.preprocessing import StandardScaler

# regression model related packages
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error

# visualization related packages
from prettytable import PrettyTable
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

# dataset loading from library
from sklearn.datasets import load_boston
```

2 Load the dataset

```
In [2]: boston_data = load_boston()
boston_data = boston_data

X = boston_data.data
y = boston_data.target
```

3 Scale the dataset

```
In [3]: scaler = StandardScaler()
X = scaler.fit_transform(X)
#y = scaler.fit_transform(y.reshape(-1, 1))
y = y.flatten()
```

```
In [4]: print('Shape of featrues %d,%d, shape of labels : %d'%(X.shape + (y.shape[0],)))
```

Shape of featrues 506,13, shape of labels : 506

4 UTIL function

```
In [5]: def plot_the_prediction(actual_values, predicted_values, title_str):
        # plot two curves
        plt.plot(range(len(actual_values)), actual_values, label='Actual')
        plt.plot(range(len(predicted_values)), predicted_values, label='Predicted')
        plt.xlabel('Iteration Number')
        plt.ylabel('Label Value')
        plt.title(title_str)
        plt.legend()
        plt.show()
```

5 sklearn implementation

```
In [6]: sk_reg = SGDRegressor()
        sk_reg.fit(X, y)
        predicted_values = sk_reg.predict(X)
        mse_sklearn = mean_squared_error(y, predicted_values)
        print('MSE for skllearn version: ', mse_sklearn)
```

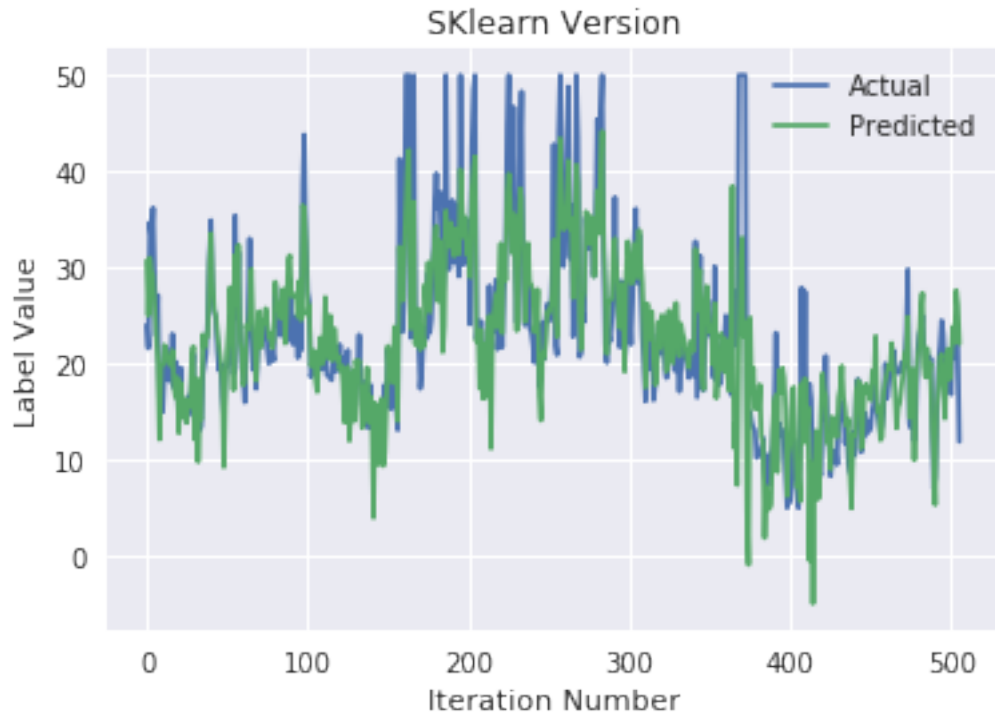
MSE for skllearn version: 22.931851381723657

```
/home/amd_3/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/stochastic_gradient.py:12
    "and default tol will be 1e-3." % type(self), FutureWarning)
```

```
In [7]: sk_weights = sk_reg.coef_
        sk_bias = sk_reg.intercept_
```

5.0.1 Y vs Y_ plot for sklearn implementation

```
In [8]: title_str = 'SKlearn Version'
        plot_the_prediction(y, predicted_values, title_str)
```



```
In [9]: comparison_df = pd.DataFrame({'Actual':y, 'Predicted':predicted_values}, index=range(len(y)))
comparison_df.head()
```

```
Out[9]:
```

	Actual	Predicted
0	24.0	30.692099
1	21.6	25.019310
2	34.7	31.014886
3	33.4	29.664846
4	36.2	29.131055

6 Custom Implementation

```
In [10]: class SGD_Linear_Regression:

    def __init__(self, with_momentum=False):

        # initialize the required variables
        self.num_iters = 60000
        self.batch_size = 64

        # variables for weight & intercept
        self.weights = np.nan
```

```

    # hyper params
    self.with_momentum = with_momentum
    self.tol = 1e-8 # quit the loop if weight chage is below this threshold
    self.momentum = 0.20
    self.learning_rate = 0.001

    # Parameters of linear regression
    self.coef_ = np.nan
    self.intercept_ = np.nan

def adjust_param(self, val, min_bound, max_bound):
    """
    This function does a linear change in momentum or learning rate as the number
    iterations increases.
    """
    mappd_value = min_bound + (max_bound - min_bound) * ((val - 1) / (self.num_iter

    return mappd_value

def fit(self, X, y):
    """
    This function fit to the data using Stochastic Gradient Descent
    """
    # initialize the weight + bias
    self.weights = np.random.uniform(low=-1.0, high=1.0, size=X.shape[1] + 1)

    # to store previous gradients for momentum
    self.prev_grad = np.zeros(X.shape[1] + 1)

    print('inital weights assigned', self.weights)

    # create a data frame for training input
    train_df = pd.DataFrame(X)
    train_df['Target'] = y

    # run multiple iterations
    for index in range(1, self.num_iters + 1):

        # sample the dataset (current batch of size self.batch_size)
        sample_df = train_df.sample(n=self.batch_size)

        # split data set to features & labels
        labels_df = sample_df['Target']
        features_df = sample_df.drop(['Target'], axis=1)

```

```

# add 1 to the last column to account bias as weight value
features_df['bias'] = 1

# compute the predicted values
predicted_values = (self.weights * features_df).sum(axis=1)

# compute loss values
loss_values = labels_df - predicted_values

# compute the error
mse_error = mean_squared_error(labels_df, predicted_values)

# print the error
if (index-1) % 10000 == 0:
    print('Iteration %d loss value: %f : lr:%f'%(index, mse_error, self.learning_rate))

# get the current gradient
gradient_weight = -2 * (features_df.mul(loss_values, axis=0)).mean(axis=0)

# run the update step for momentum
if self.with_momentum:

    # momentum term
    gradient_weight = self.momentum * self.prev_grad + self.learning_rate * gradient_weight
    self.prev_grad = gradient_weight
    self.momentum = self.adjust_param(index, 0.10, 0.90)

else: # non-momentum version
    gradient_weight *= self.learning_rate

# compute new weight vector
new_weights = self.weights - gradient_weight

# adjust learning rate
self.learning_rate = 1e-3 - self.adjust_param(index, 1e-8, 1e-3)

# terminate early if already converged
if np.linalg.norm(new_weights - self.weights) < self.tol:
    break

# update current weight as new weight
self.weights = new_weights

# Assign the final weights and biases
self.coef_ = self.weights[:-1]
self.intercept_ = np.array(self.weights[-1:])

def predict(self, X):

```

```

        # convert X to data f
        return np.dot(X, self.coef_) + self.intercept_

```

```

In [11]: sgd_reg_without_momentum = SGD_Linear_Regression(with_momentum=False)
        sgd_reg_with_momentum = SGD_Linear_Regression(with_momentum=True)

```

6.1 a) Without momentum version

```

In [12]: sgd_reg_without_momentum.fit(X, y)
        predicted_values = sgd_reg_without_momentum.predict(X)
        mse_without_momentum = mean_squared_error(y, predicted_values)
        print('The final MSE value without momentum', mse_without_momentum)
        # get bias and weights
        without_weights = sgd_reg_without_momentum.coef_
        withouts_bias = sgd_reg_without_momentum.intercept_

```

```

initial weights assigned [-0.10699758  0.1326927 -0.59571568 -0.74806401 -0.61674516 -0.83652181
 0.22077201  0.53785738  0.24364007 -0.50370687 -0.63793894  0.30277594
-0.29353685 -0.52790142]

```

```

Iteration 1 loss value: 624.118738 : lr:0.001000

```

```

Iteration 10001 loss value: 20.217794 : lr:0.000833

```

```

Iteration 20001 loss value: 14.649243 : lr:0.000667

```

```

Iteration 30001 loss value: 23.965127 : lr:0.000500

```

```

Iteration 40001 loss value: 19.374366 : lr:0.000333

```

```

Iteration 50001 loss value: 12.894435 : lr:0.000167

```

```

The final MSE value without momentum 21.89824079416991


```

6.1.1 Y vs Y_ plot for custom version without momentum

```

In [13]: title_str = 'custom version without momentum'
        plot_the_prediction(y, predicted_values, title_str)

```



```
In [14]: comparison_df = pd.DataFrame({'Actual':y, 'Predicted':predicted_values}, index=range(len(y)))
comparison_df.head()
```

```
Out[14]:
```

	Actual	Predicted
0	24.0	30.031080
1	21.6	25.031326
2	34.7	30.571907
3	33.4	28.612655
4	36.2	27.951806

6.2 b) With momentum version

```
In [15]: sgd_reg_with_momentum.fit(X, y)
predicted_values = sgd_reg_with_momentum.predict(X)
mse_withmomentum = mean_squared_error(y, predicted_values)
print('The final MSE value with momentum', mse_withmomentum)
# get bias and weights
with_weights = sgd_reg_with_momentum.coef_
with_bias = sgd_reg_with_momentum.intercept_
```

```
initial weights assigned [ 0.34436705  0.05880664 -0.58824649  0.7545802   0.66926351  0.46999522
 -0.26835904 -0.85713751 -0.57609302 -0.00371549  0.85743946  0.03037301
 -0.27145042 -0.21153613]
```

```
Iteration 1 loss value: 527.617726 : lr:0.001000
```

```

Iteration 10001 loss value: 26.383148 : lr:0.000833
Iteration 20001 loss value: 19.388833 : lr:0.000667
Iteration 30001 loss value: 16.247267 : lr:0.000500
Iteration 40001 loss value: 18.176866 : lr:0.000333
Iteration 50001 loss value: 20.411124 : lr:0.000167
The final MSE value with momentum 21.898794101026567

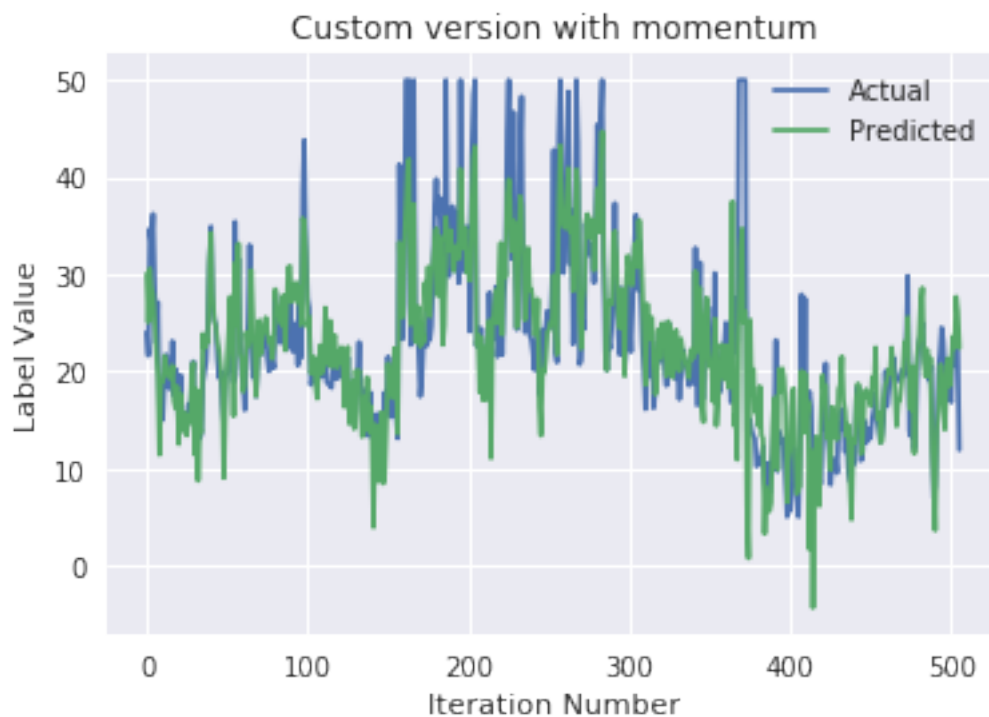
```

6.2.1 Y vs Y_ plot for Custom version with momentum

```

In [16]: title_str = 'Custom version with momentum'
         plot_the_prediction(y, predicted_values, title_str)

```



```

In [17]: comparison_df = pd.DataFrame({'Actual':y, 'Predicted':predicted_values}, index=range(len(y)))
         comparison_df.head()

```

```

Out[17]:   Actual  Predicted
0     24.0    30.016968
1     21.6    25.032358
2     34.7    30.575248
3     33.4    28.623558
4     36.2    27.942372

```


6.3 Comparision of Weights from all three models

```
In [18]: weights_table = pd.DataFrame({'SK_Learn':sk_weights, 'Without_momentum':without_weights,
                                       'With_momentum' : with_weights}, index=range(len(sk_weights)))
        bias_table = pd.DataFrame({'SK_Learn':sk_bias, 'Without_momentum': withouts_bias,
                                   'With_momentum' : with_bias})
```

6.4 a) Weights table

```
In [19]: weights_table = weights_table[['SK_Learn', 'Without_momentum', 'With_momentum']]
        weights_table
```

```
Out [19]:
```

	SK_Learn	Without_momentum	With_momentum
0	-0.752903	-0.906675	-0.922864
1	0.597092	1.074215	1.091047
2	-0.508069	0.124804	0.143693
3	0.775476	0.685810	0.690379
4	-1.044620	-2.051741	-2.067917
5	3.147765	2.676833	2.650176
6	-0.261840	0.022977	0.021907
7	-2.043230	-3.096169	-3.118597
8	0.844388	2.610523	2.664484
9	-0.524075	-2.034974	-2.079862
10	-1.803282	-2.060391	-2.048109
11	0.915453	0.850923	0.861836
12	-3.412783	-3.740848	-3.784314

6.5 b) Bias Table

```
In [20]: bias_table = bias_table[['SK_Learn', 'Without_momentum', 'With_momentum']]
        bias_table
```

```
Out [20]:
```

	SK_Learn	Without_momentum	With_momentum
0	22.341985	22.531998	22.533082

6.6 Observation

Bias holds larger values when compared with weights

7 Predict House Price given a Raw Data Point

7.1 Proceudre for predicting the house price

As the first step, we need to scale the given data point (x) to scaled domain using the parameters of standard scaler (the scaler we used at the time of training)

We need to use two parameters mean (μ) and standard deviation (σ) of the scaler and scale it using $\frac{x-\mu}{\sigma}$

Once the scaling is done, we can feed the point to model and get the predicted value as output from it

```
In [21]: def predict_the_price(model, scaler, x):
        """
        This function helps to predict the price of a data point
        given the model and its scaler.
        """

        # get parameters
        mu = scaler.mean_
        sigma = np.sqrt(scaler.var_)

        # scale the input data
        x_new = (x - mu) / sigma

        # predict the price
        predicted_price = model.predict(x_new)

        return predicted_price
```

7.2 Raw input data

```
In [22]: raw_data_points = boston_data.data[0:5]
        raw_data_points_labels = boston_data.target[0:5]
```

7.3 Predict the price

7.3.1 a) using the sklearn version

```
In [23]: predicted_values_ver1 = predict_the_price(sk_reg, scaler, raw_data_points)
```

7.3.2 b) using the custom without momentum version

```
In [24]: predicted_values_ver2 = predict_the_price(sgd_reg_without_momentum, scaler, raw_data_points)
```

7.3.3 c) using the custom with momentum version

```
In [25]: predicted_values_ver3 = predict_the_price(sgd_reg_with_momentum, scaler, raw_data_points)
```

7.4 Prediction Results

```
In [26]: predicted_df = pd.DataFrame({'SK_learn' : predicted_values_ver1,
        'Custom_without_momentum' : predicted_values_ver2,
        'Custom_with_momentum' : predicted_values_ver3,
        'Actual' : raw_data_points_labels}, index=range(len(predicted_values_ver1))
predicted_df = predicted_df[['SK_learn', 'Custom_without_momentum', 'Custom_with_momentum', 'Actual']]
predicted_df
```

```
Out[26]:
```

	SK_learn	Custom_without_momentum	Custom_with_momentum	Actual
0	30.692099	30.031080	30.016968	24.0
1	25.019310	25.031326	25.032358	21.6

2	31.014886	30.571907	30.575248	34.7
3	29.664846	28.612655	28.623558	33.4
4	29.131055	27.951806	27.942372	36.2

8 Procedure Summary

Scale the features of dataset using standard scaler

Train all the three models using the scaled dataset

Evaluate the models using MSE metric

Compare the weights, biases obtained for all models

Compare the actual and predicted values of all models

Predict the house price given a non-scaled raw input data set

9 Results Summary

```
In [27]: from prettytable import PrettyTable
```

```
In [28]: ptable = PrettyTable()
         ptable.title = 'SGD Regression Method Comparison'
         ptable.field_names = ['Method', 'MSE Value']
```

```
In [29]: ptable.add_row(['Sk-learn Version', mse_sklearn])
         ptable.add_row(['Custom Version without momentum', mse_without_momentum])
         ptable.add_row(['Custom Version with momentum', mse_withmomentum])
```

```
In [30]: print(ptable)
```

```
+-----+
|          SGD Regression Method Comparison          |
+-----+-----+
|          Method          |          MSE Value          |
+-----+-----+
|          Sk-learn Version          | 22.931851381723657 |
| Custom Version without momentum | 21.89824079416991 |
| Custom Version with momentum | 21.898794101026567 |
+-----+-----+
```

10 Conclusion

The weights obtained from both momentum and non momentum version are very similar

Many of the weights obtained from sklearn implementation are close to the custom implementation weights

The final MSE error from all three methods are roughly same

The performace of momentum version is better compared to non-momentum version (from the mse values obtained at different iterations 10K, 20K, 30K etc.)